22-Jan 2025

CROSS VALIDATION

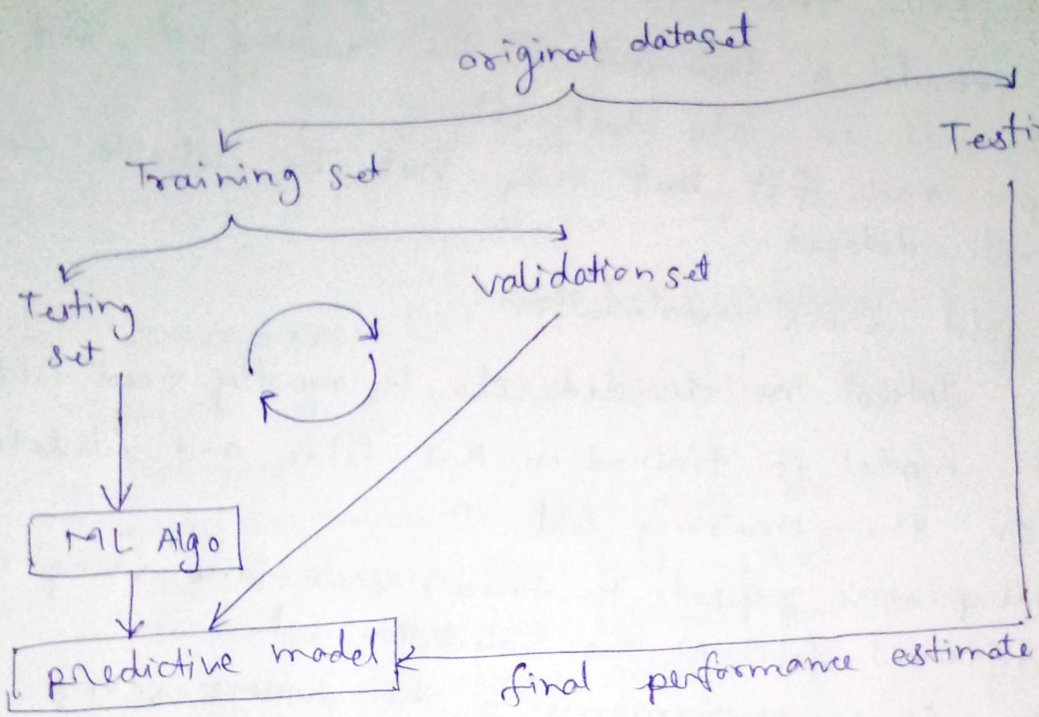

Cross validation:

- Cross validation is a resampling technique used to assess the performance and generalization ability of a ML model.
- it helps in mitigating overfitting and provides a better estimate of how the model will perform on unseen data.
- instead of using a single train-test split, cross-validation divides the dataset into multiple subsets, training the model on some and validating it on others in a systematic manner.

Types of cross validation.

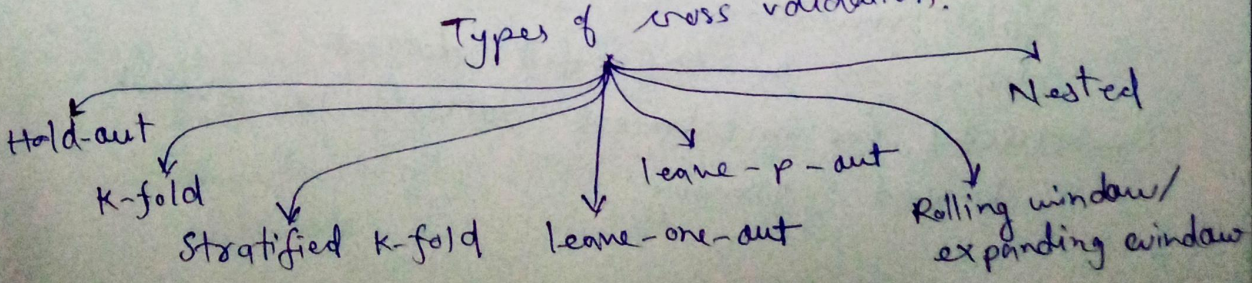

(i) Hold-Out Cross-Validation:

- The dataset is split into training and test sets (e.g. 80% training, 20% testing)
- The model is trained on the training set and evaluated on the test set.
- Simple and fast but may not be reliable for small dataset.

(ii) K-fold cross validation:

- The dataset is divided into k equally sized folds.
- The model is trained on k-1 folds and validated on the remaining fold.
- This process repeats k-times, each time using a different fold as the validation set.
- The final performance is the average score all k iterations.

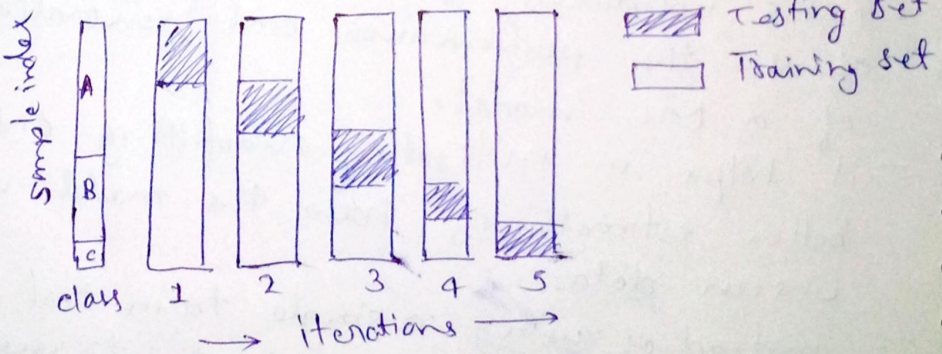

(iii) Stratified k-fold Cross validation:

- similar to k-fold CV, but ensures that each fold maintain the same proportion of target classes as in the full dataset.
- useful for imbalanced datasets to prevent biased training.

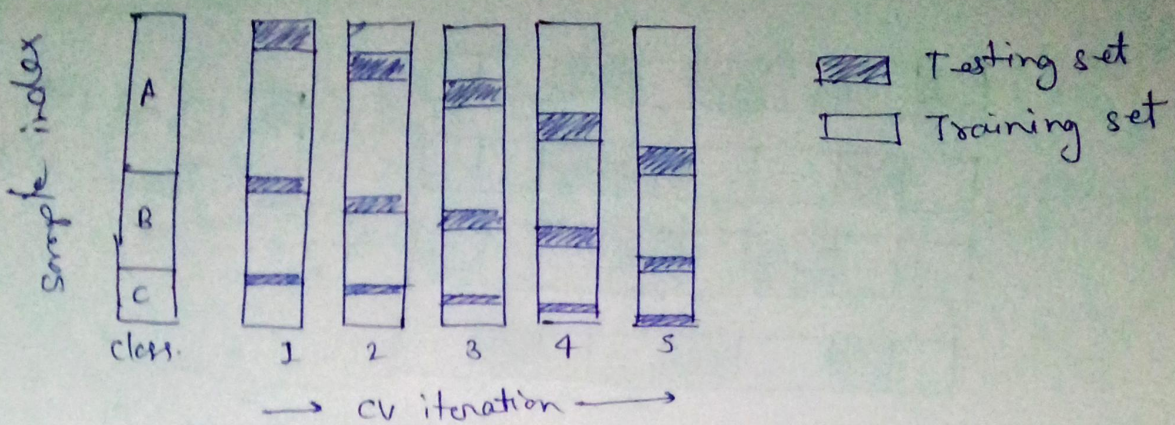

(iv) leave-one-Out Cross-Validation (LOOCV):

- special case of k-fold CV where k = number of samples
- each instance is used once as a test set while the rest form the training set.
- Gives an unbiased estimate but is computationally expensive for large datasets.

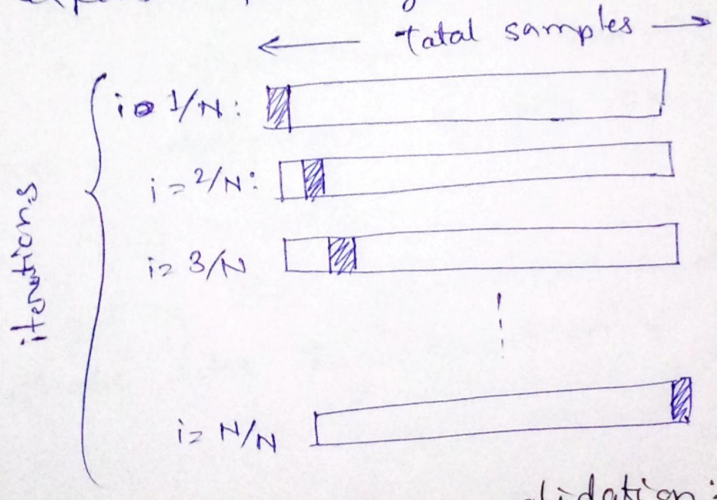

(v) leave-p-out Cross validation:

- similar to LOOCV but instead of 1, p samples are left out for validation
- Repeated multiple times for all possible subsets of size p.
- computationally expensive but provides robust validation.

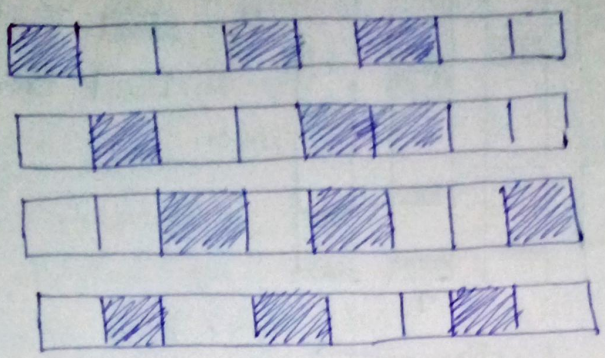

Test set

Train set